

Adversarial Robustness of U-Net for Inland Water Body Segmentation from SAR Images

Siddharth Kothari
siddharth.kothari@iiitb.ac.in
Graphics-Visualization-Computing
Lab, International Institute of
Information Technology, Bangalore
Bangalore, Karnataka, India

M Srinivasan
M.Srinivasan@iiitb.ac.in
Graphics-Visualization-Computing
Lab, International Institute of
Information Technology, Bangalore
Bangalore, Karnataka, India

Sankalp Kothari
sankalp.kothari@iiitb.ac.in
Graphics-Visualization-Computing
Lab, International Institute of
Information Technology, Bangalore
Bangalore, Karnataka, India

Ujjwal Verma
ujjwal.verma@manipal.edu
Department of Electronics and
Communication Engineering,
Manipal Institute of Technology,
Manipal Academy of Higher
Education
Manipal, Karnataka, India

Jaya Sreevalsan-Nair
jnair@iiitb.ac.in
Graphics-Visualization-Computing
Lab, International Institute of
Information Technology, Bangalore
Bangalore, Karnataka, India

ABSTRACT

Inland water body segmentation from Synthetic Aperture Radar (SAR) images is crucial for applications such as flood monitoring and environmental assessment. While convolutional neural networks (CNNs), particularly U-Net, achieve state-of-the-art performance in SAR segmentation, they depend heavily on manually annotated masks for training. These annotations are prone to label noise, especially around complex river geometries, which can propagate to the models. In this work, we simulate annotation errors using morphological label noise—generated via controlled erosion and dilation, and evaluate the robustness of U-Net under such data poisoning attacks. Our findings show that U-Net can tolerate a limited degree of corruption before performance degrades sharply, underscoring the importance of annotation quality. We also release the dataset, code, and adversarial examples to facilitate robust model development. (GitHub link: <https://github.com/GVCL/IWSeg-SAR-Poison.git>)

KEYWORDS

Deep learning, SAR Images, Image Segmentation, U-Net, Label Noise, Adversarial Attacks, Morphological Operations

ACM Reference Format:

Siddharth Kothari, M Srinivasan, Sankalp Kothari, Ujjwal Verma, and Jaya Sreevalsan-Nair. 2024. Adversarial Robustness of U-Net for Inland Water Body Segmentation from SAR Images. In . ACM, New York, NY, USA, 3 pages. <https://doi.org/10.1145/3681769.3698585>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference'17, July 2017, Washington, DC, USA

© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 979-8-4007-1148-0/24/10
<https://doi.org/10.1145/3681769.3698585>

1 INTRODUCTION




Figure 1: Example of Sentinel-1 SAR images of the Amazon River Basin (left) in our dataset, along with their generated binary water masks (right).

Water body segmentation is vital in remote sensing, supporting flood mapping [5, 7, 10], glacial lake monitoring [11], sedimentation studies [6], and environmental assessments. Synthetic Aperture Radar (SAR) offers all-weather, high-resolution data, yet inland water segmentation remains difficult due to the complex geometry of river basins and tributaries.

Deep learning methods, particularly encoder-decoder CNNs like U-Net [9], surpass traditional approaches but depend heavily on annotated masks, which often contain errors from boundary ambiguities, width variations, and mislabeling [2, 12]. Such errors can impair training, motivating us to study their impact through adversarial robustness.

Unlike typical adversarial attacks based on imperceptible perturbations [1, 8], we introduce *morphological label noise*—annotation corruptions via erosion and dilation—as a proxy for human errors in complex river geometries (e.g., the Amazon Basin). We evaluate U-Net under varying corruption levels to assess its tolerance to noisy supervision.

Our main contributions are:

- (1) A framework to simulate annotation errors with morphological operations, and a dataset with adversarial examples.
- (2) An analysis of U-Net's robustness across corruption levels.

2 METHOD

We generated our dataset using Sentinel-1 SAR imagery for the Amazon Basin, with ground truth water masks derived from Sentinel-2 NDWI thresholding. An example SAR image and mask is shown in Figure 1, and the detailed dataset preparation process can be found in [4].

2.1 Simulation of Adversarial Attacks in Training

Erosion and dilation are morphological operations commonly used in image processing, which we apply to generate controlled corruption of annotated masks. Since human annotators may introduce errors at land–water boundaries, especially in complex hydrological geometries, these operations are well-suited as they add or remove pixels at class boundaries. To simulate real-world non-uniformity, we introduce randomness in the kernel sizes during corruption. For images of size 512×512 , we apply corruption at levels $\{2\%, 12\%, 15\%, 17\%, 20\%, 25\%, 30\%\}$. The resulting corrupted masks are then used in adversarial experiments to assess model robustness against human-like boundary errors.

- (1) Out of the train and the validation images, we select one-third of the images for corruption using erosion, one-third using dilation, and one-third is left uncorrupted.
- (2) For an experiment, we then set a threshold on the number of corrupted pixels allowed.
- (3) Erosion and dilation are applied to the corresponding masks iteratively, till either a pre-defined number of iterations are over, or the number of corrupted pixels crosses the threshold, following which we choose the previous highest corrupted version.

Although we set limits on the number of corrupted pixels, these are not always reached in practice, particularly for erosion. Since erosion reduces white pixels, it saturates once an insufficient number remain, preventing further corruption. This effect is illustrated in Figure 2, which shows box plots of corrupted pixels for erosion and dilation across all training images. For erosion, the number of corrupted pixels stagnates beyond 20%, with median and lower quartile values remaining almost unchanged, while dilation follows the expected increasing trend.

3 EXPERIMENTS & RESULTS

Table 1 presents the evaluation metrics of the U-Net model across different corruption levels. The model achieves around 95% segmentation accuracy, with performance largely unaffected up to 20% corruption and only moderate degradation beyond 25%. Figure 3 further illustrates this effect: as erosion increases, the proportion of white pixels decreases, yet the U-Net output, though blurred before thresholding, still captures the river geometry accurately. Thresholding restores clear masks, confirming the model's robustness to substantial corruption.




Figure 2: A box plot showing the range of actual corruption in the training data for different corruption levels.

Table 1: Performance of the baseline U-Net model on our dataset at various corruption levels

Corruption Level	Dice Coeff.	IoU	Precision	Recall	F1 Score	Specificity	Accuracy
0%	0.770	0.823	0.896	0.869	0.865	0.971	0.955
2%	0.775	0.825	0.896	0.873	0.868	0.968	0.955
12%	0.769	0.822	0.899	0.864	0.865	0.975	0.955
15%	0.768	0.819	0.890	0.866	0.863	0.971	0.954
17%	0.758	0.814	0.903	0.854	0.859	0.979	0.952
20%	0.725	0.790	0.902	0.830	0.841	0.981	0.946
25%	0.597	0.708	0.875	0.759	0.768	0.984	0.920
30%	0.253	0.496	0.818	0.588	0.567	0.996	0.828

We selected U-Net for adversarial experimentation due to its strong performance with smaller training sets, faster training, and stability to annotation errors [3]. Future work includes extending experiments to other segmentation architectures, analyzing erosion and dilation independently to quantify their impact, and exploring traditional corruptions such as rotation and translation.




Figure 3: Qualitative results showing the effect of corruption level on the ground truth (second column from the left) and the final predicted (rightmost column) mask, for the input image shown in Figure 1 (leftmost column). The predicted masks before applying the threshold (second column from the right) to the U-Net model outcomes get blurrier with an increase in corruption, but are robust in capturing river structure in the presence of morphological label noise.

REFERENCES

- [1] A. Edpuganti, P. Akshaya, J. Gouthami, V. V. Sajith Variyar, V. Sowmya, and R. Sivanpillai. 2023. Effect of Data Quality on Water Body Segmentation with DeepLabV3+ Algorithm. *The International Archives of the Photogrammetry, Remote Sensing and Spatial Information Sciences* XLVIII-M-3-2023 (2023), 81–85. <https://doi.org/10.5194/isprs-archives-XLVIII-M-3-2023-81-2023>
- [2] Seyed Mohammad Hassan Erfani, Zhenyao Wu, Xinyi Wu, Song Wang, and Erfan Goharian. 2022. ATLANTIS: A Benchmark for Semantic Segmentation of Waterbody Images. *Environmental Modelling & Software* 149 (2022), 105333.
- [3] Nicholas Heller, Joshua Dean, and Nikolaos Papanikolopoulos. 2018. Imperfect Segmentation Labels: How much do they matter?. In *Intravascular Imaging and Computer Assisted Stenting and Large-Scale Annotation of Biomedical Data and Expert Label Synthesis: 7th Joint International Workshop, CVII-STENT 2018 and Third International Workshop, LABELS 2018, Held in Conjunction with MICCAI 2018, Granada, Spain, September 16, 2018, Proceedings* 3. Springer, 112–120.
- [4] Siddharth Kothari, Srinivasan Murali, Sankalp Kothari, Ujjwal Verma, and Jaya Sreevalsan-Nair. 2025. Adversarial Robustness of Deep Learning Models for Inland Water Body Segmentation from SAR Images. arXiv:2505.01884 [eess.IV] <https://arxiv.org/abs/2505.01884>
- [5] Laura Lopez-Fuentes, Claudio Rossi, and Harald Skinnemoen. 2017. River Segmentation for Flood Monitoring. In *2017 IEEE International Conference on Big Data (Big Data)*. 3746–3749. <https://doi.org/10.1109/BigData.2017.8258373>
- [6] Conor O'Sullivan, Ambrish Kashyap, Seamus Coveney, Xavier Monteys, and Soumyabrata Dev. 2024. Enhancing Coastal Water Body Segmentation with Landsat Irish Coastal Segmentation (LICS) Dataset. *Remote Sensing Applications: Society and Environment* 36 (2024), 101276. <https://doi.org/10.1016/j.rsase.2024.101276>
- [7] Yanping Qin, Xiaobin Yin, Yan Li, Qing Xu, Lei Zhang, Peng Mao, and Xingwei Jiang. 2025. High-Precision Flood Mapping From Sentinel-1 Dual-Polarization SAR Data. *IEEE Transactions on Geoscience and Remote Sensing* 63 (2025), 1–15. <https://doi.org/10.1109/TGRS.2025.3557330>
- [8] Mustafizur Rahaman, Md. Monsur Hillas, Jannatul Tuba, Jannatul Ferdous Ruma, Nahian Ahmed, and Rashedur M. Rahman and. 2022. Effects of Label Noise on Performance of Remote Sensing and Deep Learning-based Water Body Segmentation Models. *Cybernetics and Systems* 53, 6 (2022), 581–606. <https://doi.org/10.1080/01969722.2021.1989171> arXiv:<https://doi.org/10.1080/01969722.2021.1989171>
- [9] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. 2015. U-Net: Convolutional Networks for Biomedical Image Segmentation. In *Proceedings of 18th International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI 2015), part III* 18. Springer, 234–241.
- [10] Tamer Saleh, Shima Holail, Mohamed Zahran, Xiongwu Xiao, and Gui-Song Xia. 2024. LiST-Net: Enhanced Flood Mapping With Lightweight SAR Transformer Network and Dimension-Wise Attention. *IEEE Transactions on Geoscience and Remote Sensing* 62 (2024), 1–17. <https://doi.org/10.1109/TGRS.2024.3397797>
- [11] Shidong Wang, Maria V. Peppas, Wen Xiao, Sudan B. Maharjan, Sharad P. Joshi, and Jon P. Mills. 2022. A Second-order Attention Network for Glacial Lake Segmentation from Remotely Sensed Imagery. *ISPRS Journal of Photogrammetry and Remote Sensing* 189 (2022), 289–301. <https://doi.org/10.1016/j.isprsjprs.2022.05.007>
- [12] Junshi Xia, Hongruixuan Chen, Clifford Broni-Bediako, Yimin Wei, Jian Song, and Naoto Yokoya. 2025. OpenEarthMap-SAR: A Benchmark Synthetic Aperture Radar Dataset for Global High-Resolution Land Cover Mapping. arXiv:2501.10891 [eess.IV] <https://arxiv.org/abs/2501.10891>